

Open your .env file and set your PostgreSQL connection string:

**DATABASE\_URL="postgresql://user:password@localhost:5432/yourdbname"**

◇ Replace user, password, and yourdbname with your actual PostgreSQL credentials.

# PostgreSQL Notes

---

## 1) Quick intro

- **PostgreSQL** = powerful open-source relational DBMS.
- Uses **SQL** (ANSI standard) plus many advanced features: **JSON/JSONB, ARRAY, UUID**, custom types, window functions, full-text search, stored procedures, etc.

## 2) Installation & tools (short)

- Windows / macOS: installers available (Postgres.app on macOS).
- Linux (Debian/Ubuntu): `sudo apt install postgresql postgresql-contrib`
- Tools: **psql** (CLI), **pgAdmin** (GUI), **DBeaver** (multi-DB GUI).

Start/stop (Linux example):

```
sudo systemctl start postgresql  
sudo systemctl enable postgresql
```

Connect with psql as default user:

```
sudo -u postgres psql  
-- or from local: psql -U youruser -d yourdb -h localhost -p 5432
```

## 3) psql meta-commands (very useful)

Run inside psql:

```
\l          -- list databases  
\c dbname   -- connect to db  
\dt        -- list tables  
\d tablename -- describe table (columns, types, constraints)  
\df        -- list functions  
\?        -- help for psql commands  
\q        -- quit
```

## 4) Create DB / User / Role / Permissions

```
-- create role (user) with password
CREATE ROLE sagar WITH LOGIN PASSWORD 'strongpass';

-- create database owned by role
CREATE DATABASE practice_db OWNER sagar;

-- grant privileges (if needed)
GRANT ALL PRIVILEGES ON DATABASE practice_db TO sagar;
```

---

## 5) Basic schema & CRUD (examples)

```
-- create table
CREATE TABLE students (
    id SERIAL PRIMARY KEY,
    name VARCHAR(100) NOT NULL,
    email TEXT UNIQUE,
    age INT,
    enrolled BOOLEAN DEFAULT TRUE,
    created_at TIMESTAMP DEFAULT NOW()
);

-- insert
INSERT INTO students (name, email, age) VALUES
('Sagar', 'sagar@example.com', 21),
('Maya', 'maya@example.com', 23);

-- select
SELECT * FROM students;
SELECT name, age FROM students WHERE age > 21 ORDER BY age DESC
LIMIT 5;

-- update
UPDATE students SET age = age + 1 WHERE id = 1;

-- delete
DELETE FROM students WHERE id = 2;
```

---

## 6) Important Data types

- `SERIAL / BIGSERIAL` — auto increment integer (backed by sequence)
- `INTEGER, BIGINT, SMALLINT`
- `VARCHAR(n), TEXT`
- `BOOLEAN`
- `TIMESTAMP, DATE, TIME`
- `NUMERIC(p,s)` — precise decimals (money)
- `UUID` — universally unique id (use `gen_random_uuid()` from `pgcrypto` or `uuid-ossp`)
- `JSON, JSONB` — JSONB recommended for indexing and performance
- `ARRAY` — e.g., `TEXT[]`
- `ENUM` — custom enum types

Example JSONB:

```
CREATE TABLE products (
    id SERIAL PRIMARY KEY,
    data JSONB
);

INSERT INTO products (data) VALUES
('{"name": "shoe", "sizes": [38, 39, 40], "stock": 10}');
SELECT data->>'name' AS name FROM products;
```

---

## 7) Constraints & Keys

- PRIMARY KEY, UNIQUE, NOT NULL, CHECK, FOREIGN KEY

```
ALTER TABLE students ADD CONSTRAINT chk_age CHECK (age >= 0);
```

Foreign key:

```
CREATE TABLE courses (
    id SERIAL PRIMARY KEY,
    student_id INT REFERENCES students(id) ON DELETE CASCADE,
    course_name VARCHAR(100)
);
```

---

## 8) Indexes & when to use

- Create index to speed up queries on columns used in WHERE/JOIN/ORDER BY:

```
CREATE INDEX idx_students_name ON students(name);
```

- Use EXPLAIN ANALYZE <query> to see query plan and whether index used.
  - Avoid useless indexes (costly on writes).
- 

## 9) Joins & subqueries

```
-- inner join
SELECT s.name, c.course_name
FROM students s
JOIN courses c ON s.id = c.student_id;

-- left join
SELECT s.name, c.course_name
```

```
FROM students s
LEFT JOIN courses c ON s.id = c.student_id;

-- subquery
SELECT name FROM students WHERE id IN (SELECT student_id FROM
courses WHERE course_name='Math');
```

---

## 10) Aggregation, GROUP BY, HAVING

```
SELECT COUNT(*) AS total, AVG(age) AS avg_age FROM students;
SELECT student_id, COUNT(*) AS course_count FROM courses GROUP BY
student_id HAVING COUNT(*) > 1;
```

---

## 11) Window functions (powerful)

```
SELECT id, name, age,
       ROW_NUMBER() OVER (ORDER BY age DESC) AS rn,
       AVG(age) OVER () AS overall_avg
FROM students;
```

---

## 12) Transactions & concurrency

- Use transactions to group operations:

```
BEGIN;
UPDATE accounts SET balance = balance - 100 WHERE id = 1;
UPDATE accounts SET balance = balance + 100 WHERE id = 2;
COMMIT;
-- or ROLLBACK to cancel
```

- Isolation levels: **READ COMMITTED** (default), **REPEATABLE READ**, **SERIALIZABLE**.  
Use when needed to avoid race conditions.

---

## 13) Stored functions & procedures

- Create functions (SQL or PL/pgSQL):

```
CREATE FUNCTION add_numbers(a INT, b INT) RETURNS INT AS $$  
BEGIN  
    RETURN a + b;  
END;  
$$ LANGUAGE plpgsql;  
  
SELECT add_numbers(2,3);
```

---

## 14) Views & materialized views

- View = virtual table

```
CREATE VIEW active_students AS SELECT id, name FROM students WHERE  
enrolled;
```

- Materialized view = stores results (refresh manually)

```
CREATE MATERIALIZED VIEW mv_popular_courses AS  
SELECT course_name, COUNT(*) cnt FROM courses GROUP BY course_name;  
REFRESH MATERIALIZED VIEW mv_popular_courses;
```

---

## 15) Full-text search (basic)

```
ALTER TABLE articles ADD COLUMN tsv tsvector;  
UPDATE articles SET tsv = to_tsvector(coalesce(title,'') || ' ' ||  
coalesce(content,''));  
CREATE INDEX idx_articles_tsv ON articles USING GIN(tsv);
```

```
SELECT * FROM articles WHERE tsv @@ to_tsquery('postgres & tutorial');
```

---

## 16) JSONB tips

- Query nested JSONB items:

```
SELECT data->'sizes' AS sizes, data->>'name' AS name FROM products;  
-- existence operator  
SELECT * FROM products WHERE data @> '{"stock":10}';
```

- Index JSONB with GIN for containment queries:

```
CREATE INDEX idx_products_data ON products USING GIN (data);
```

---

## 17) Backup & restore

- Backup:

```
pg_dump -U sagar -d practice_db > practice_db.sql  
-- for whole cluster: pg_dumpall > alldb.sql
```

- Restore:

```
psql -U sagar -d practice_db < practice_db.sql
```

- Binary backup via `pg_basebackup` for replication.
-

## 18) Roles & security best practices

- Use roles with least privileges.
  - Don't use superuser for app connections.
  - Use `pg_hba.conf` to configure authentication (md5, scram-sha-256).
  - Keep backups, use SSL for remote connections.
- 

## 19) Useful built-in functions & operators (cheat sheet)

- Date/time: `NOW()`, `CURRENT_DATE`, `AGE(timestamp)`
  - String: `LENGTH()`, `UPPER()`, `LOWER()`, `CONCAT()`
  - Aggregates: `COUNT()`, `SUM()`, `AVG()`, `MIN()`, `MAX()`
  - JSON: `->`, `->>`, `#>` (path), `@>` (contains)
  - Array: `array_length`, `unnest()`
  - Sequence: `nextval('seqname')`, `currval()`
- 

## 20) Performance tips

- Index columns used frequently in WHERE/JOIN.
- Use `EXPLAIN ANALYZE` to understand slow queries.
- Be cautious with `SELECT *` on large tables.
- Use connection pooling (PgBouncer) for web apps.

- Vacuum and analyze: `VACUUM` / `VACUUM ANALYZE` to maintain stats (or enable autovacuum).
- 

## 21) Example mini schema + queries (put it all together)

```
-- create users and posts
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    username VARCHAR(50) UNIQUE NOT NULL,
    joined TIMESTAMP DEFAULT NOW()
);

CREATE TABLE posts (
    id SERIAL PRIMARY KEY,
    user_id INT REFERENCES users(id),
    title TEXT,
    body TEXT,
    tags TEXT[],
    created_at TIMESTAMP DEFAULT NOW()
);

-- insert sample
INSERT INTO users (username) VALUES ('sagar'), ('maya');
INSERT INTO posts (user_id, title, tags) VALUES (1, 'Intro to PG',
ARRAY['postgres', 'sql']), (2, 'JS Tips', ARRAY['javascript']);

-- query: top posters
SELECT u.username, COUNT(p.id) AS posts
FROM users u
LEFT JOIN posts p ON u.id = p.user_id
GROUP BY u.username ORDER BY posts DESC;
```

---

## 22) Common pitfalls & notes

- `SERIAL` creates a sequence — if you restore and insert explicit ids, sequence may need `setval()` to avoid conflicts.
  - `TEXT` vs `VARCHAR(n)` — `TEXT` has no length limit. `VARCHAR(n)` enforces max length.
  - Case sensitivity: unquoted identifiers are folded to lowercase. Use double quotes `"MyTable"` to preserve case (avoid this unless needed).
  - JSONB is good, but if data has relational nature prefer normalized tables.
- 

## 23) Learning practice suggestions

- Recreate the example mini schema and practice queries.
  - Use `EXPLAIN ANALYZE` on slow queries.
  - Build a small CRUD app (Node.js + `pg` or Python + `psycopg2`) to practice connections and transactions.
  - Try one feature per day: JSONB, window functions, materialized views, etc.
- 

## 24) Quick resources (no links here — search for these)

- Official PostgreSQL docs (search "PostgreSQL documentation") — excellent, definitive.
  - Tutorials: "Postgres tutorial" / "psql cheat sheet" / "Postgres JSONB examples".
  - Practice: use pgAdmin or DBeaver and a toy database.
-